ChatGPT

AI-assisted consultation filtering and the limits of automated decision-making under EU law

In an administrative public consultation (for example, input on a power plant concession), the use of AI to summarize or filter submissions raises difficult questions under EU data protection and administrative law. In particular, GDPR Article 22 – the right not to be subject to certain automated decisions – must be interpreted alongside sectoral participation rights and general principles of good administration. Below, we examine six unresolved issues that test the boundaries of "solely" automated decisions, the definition of a "decision" itself, and the required safeguards when algorithms inform or perform administrative acts.

1. "Solely automated" decisions and superficial human review

Legal standard – meaningful human involvement: Article 22(1) GDPR prohibits decisions **"based solely on automated processing"** that have legal or similarly significant effects. The European Data Protection Board's guidance makes clear this means **no human involvement** in reaching the decision [1]. Critically, controllers **"cannot avoid the Article 22 provisions by fabricating human involvement"**: if a human merely **"routinely applies automatically generated profiles"** to individuals **without any actual influence on the result**, the decision is still considered solely automated [2]. To break the "solely automated" characterization, any human intervention must be **genuine and substantive**. The oversight should be **"meaningful, rather than just a token gesture,"** performed by someone with **"the authority and competence to change the decision"** and who **considers all the relevant data** afresh [3]. In other words, the human must effectively **reassess** the machine's proposal, not just rubber-stamp it.

Unclear threshold for meaningful review: The law does not quantify how intensive or time-consuming the review must be, but practical constraints can render a purported human review meaningless. Data protection authorities have emphasized factors like the human reviewer's expertise, independence, and the time and resources available to them [4] [5]. For example, if one person is tasked with "oversight" of thousands of AI-screened consultation submissions with only seconds to spend on each, this calls into question whether they can truly exercise judgment. The Spanish DPA cautions that if a reviewer must handle an overwhelming volume – say reading **100 reports of 100 pages each per day (implying an impossible 21 pages per minute)** – then **"in general, that would not be feasible,"** and human supervision becomes **"de facto unfeasible."** [5] In such scenarios, the nominal human involvement is arguably illusory. Indeed, German authorities interpreting Article 22 have suggested that to avoid a decision being "solely automated," the human must undertake a **"new assessment"** of the case, not perform a purely routine endorsement of the algorithm's output. In sum, while no bright-line rule exists, **minimal, perfunctory review – especially under conditions that make careful consideration impossible – will likely be insufficient** to remove a decision from Article 22's scope. It remains an open question at what point human intervention becomes substantive enough; in practice, regulators and courts will examine whether the human's role actually affects the outcome or whether the process is, in effect, automation with a human rubber stamp.

2. Algorithmic information gatekeeping as an Article 22 “decision”

Decisions via determining influence: If an AI system filters or prioritizes what information is presented to a human decision-maker, we must ask whether that filtering constitutes an automated “decision” about individuals. GDPR Article 22 applies to any **“decision”** based solely on automation that significantly affects a person. Notably, the concept of a “decision” here is broad – it **“is capable of including a number of acts which may affect the data subject in many ways,”** even intermediate steps like generating a score [6] [7]. The Court of Justice of the EU’s recent *SCHUFA* ruling illustrates this point. *SCHUFA*, a credit bureau, argued it only provided an automated credit score as a **preparatory act** and that the bank’s loan decision was made by a human. The CJEU rejected this argument and held that *SCHUFA’s* automated scoring itself qualified as **automated decision-making** subject to Article 22 [8] [9]. The key factor was that the score **played a determining role** in the final outcome: when a low credit score almost invariably led the bank to refuse a loan, the act of producing that score **“must be qualified in itself as a decision”** producing significant effects on the person [10]. In the Court’s eyes, it did not matter that a separate entity (the bank) formally made the final call – the automated profile effectively **decided the data subject’s fate**, triggering Article 22 protection.

Applied to algorithmic filtering of consultation submissions, the same principle suggests that if an AI’s exclusion of a submission **decisively shapes the ultimate decision**, it can count as an automated decision about the author of that submission. The filtering algorithm is essentially controlling what the human decision-maker knows and thereby determining which voices influence the outcome. EU data protection authorities have echoed this: an automated process that **“in a decisive way”** informs a later decision about a person is covered by Article 22 [11]. For instance, if an AI system categorizes a particular public comment as irrelevant or spam and as a result the official never considers it, the AI has effectively made a decision **adverse to the commenter’s participation**. The situation parallels *SCHUFA*: the automated output has a **predominant weight** in what happens next.

Human in the loop – genuine or not? The complication is when the final decision is formally made by a human. One line of reasoning, seen in a pre-*SCHUFA* German case, was that if humans retain discretion, the process isn’t “solely automated.” In a 2020 case on Amazon’s warehouse management algorithms, the Hannover Administrative Court found Article 22 inapplicable on grounds that the software’s recommendations did not bind human managers – there was **some human involvement**, breaking the “solely automated” criterion [12]. However, *SCHUFA* calls into question overly formalistic approaches. The CJEU’s focus was on real-world influence rather than technical decision allocation. If in practice the human decision-maker **relies heavily on the algorithm’s output** and rarely deviates from it, the filtering tool can be seen as having a “determining” or **de facto decisive** role. In short, controlling the information flow to a human can be functionally equivalent to deciding on that human’s behalf. No EU court has yet applied *SCHUFA’s* logic to an administrative consultation context, but the implication is that **upstream AI filtering can trigger Article 22** when the downstream human’s decision is a foregone conclusion shaped by the algorithm. The degree of **dependence** on the AI will be crucial. If the AI merely suggests or highlights some inputs but the human routinely reviews *all* submissions substantively, Article 22 might not be engaged. But if the AI effectively **screens out certain public inputs from any meaningful consideration**, that filtering is likely to be deemed an automated decision vis-à-vis those excluded contributors.

3. Indirect effects on third parties: does exclusion from a process "significantly affect" the person?

Article 22(1) covers automated decisions with **"legal effects"** on a person or other **"similarly significant"** effects. Here the person submitting a consultation comment is not the direct subject of the final administrative decision (which concerns the permit applicant), raising the question: can exclusion of their input count as a significant effect on them? This is a novel issue with no directly on-point case law.

On one hand, one might argue that a person has no absolute right to have their viewpoint determine a public decision, so filtering out a comment is merely a disappointment or political frustration, not a **protected individual effect**. The GDPR's examples of significant effects (refusal of credit, denial of employment, etc. [13] [14]) involve impacts on one's personal legal or economic situation. By contrast, being left out of a consultation might be characterized as a step removed: the outcome (granting the concession) is about a third party, and the commenter's position doesn't change in a legal register as a direct result.

On the other hand, public participation rights in certain procedures are themselves legally recognized interests. In environmental decision-making, for example, members of the public – especially those affected by the project – have enforceable rights under EU law to submit comments and have them taken into account. The Aarhus Convention (to which the EU is a party) requires that **"due account is taken of the outcome of the public participation"** in the final decision [15]. If an algorithmic system prevents authorities from taking a person's input into account, it arguably **deprives that person of a procedural right**. The effect is not merely abstract: it undermines the individual's ability to protect their concrete interests (such as environmental or property interests that the consultation is meant to safeguard). In that sense, being algorithmically silenced could be seen as **"similarly significant"** detriment – a procedural harm leading to a potentially substantive outcome (for instance, a plant is approved that might have been halted or modified had the comment been considered).

This issue remains unresolved. Data protection regulators have not squarely addressed whether **exclusion from influencing a decision** can qualify as a significant effect under Article 22. The conservative view is that Article 22 is intended to cover decisions *about* an individual (determining *their* rights, status, opportunities, etc.), not every instance where someone's input is ignored. The counter-argument is that modern notions of rights recognize participatory harms – especially where an individual has a legal entitlement to be heard. If the consultation in question is a legal condition of a fair administrative process, the automated filtering of a person's submission could be seen as **denial of an individual right**, producing a legal effect for the data subject (the loss of an opportunity to be heard, which could itself be actionable under administrative law). Until we see litigation or guidance on this, a cautious approach for controllers would be to treat such filtering as engaging data subject rights: at minimum, **transparency** to the commenter and an opportunity to contest the exclusion would align with the spirit of Article 22, even if its applicability is debatable. In any event, the **significance** criterion in Article 22(1) would likely be examined in light of the context – e.g. a public consultation on environmental permits (where participation is a legal right) might tilt toward recognizing a significant effect, whereas filtering comments on a non-binding online survey might not.

4. Automated procedural actions vs. decisions: routing, triage and other “housekeeping” steps

Modern algorithmic systems may not issue final decisions outright, but they increasingly perform **gatekeeping and triage functions**: routing cases to particular handlers, prioritizing some matters over others, or categorizing submissions in ways that influence how they are processed. The legal ambiguity is whether and when these automated intermediate steps themselves amount to **decisions “concerning” the individual** under Article 22. The GDPR does not explicitly distinguish between a “preparatory” act and a decision – what matters is whether an act **decides something affecting the individual’s position**. In multi-stage processes, however, identifying the decision point is challenging. Scholars have noted the **“difficulty of locating a ‘decision’ in instances of multi-stage decision-making”**, especially where there are **varying degrees of human involvement throughout the process** [16]. For example, consider an AI system that initially categorizes incoming complaints as high or low priority. A human might later decide whether to uphold or reject the complaint, but only after the low-priority ones have been shunted aside (perhaps effectively never acted upon). Is the initial categorization an automated decision in its own right?

The *SCHUFA* judgment, though dealing with a credit score, signals a broad approach: even a **“preparatory”** automated act can be a decision if it **substantively shapes the outcome** [8] [9]. The CJEU looked past formal roles and examined functional impact. Applying that logic, an automated categorization that steers a case down a path virtually guaranteeing a certain result (e.g. routing a citizen’s complaint to a “routine dismissal” queue vs. an investigative unit) could be seen as an Article 22 decision. The more the algorithm’s classification **determines or heavily dictates** the eventual handling, the stronger the argument that it is making a decision about the person. Conversely, if the automated step is truly neutral or easily corrected by human judgment (for instance, a mere sorting that does not impede any case from being fully heard in due course), it might be considered a non-decisive preparatory act.

No clear legal test exists yet. National courts have so far only nibbled at the issue. In the Amazon warehouse case mentioned above, the court was reluctant to find an Article 22 decision in an algorithm assigning work-tasks, reasoning that management could override those assignments (i.e. the tool was advisory) [12]. That reasoning underscores the importance of whether human intervention is **genuine and routine** or merely theoretical. After *SCHUFA*, one might expect a more stringent review: if in practice the **default outputs of an algorithm are followed in the vast majority of cases**, the fact that a human *could* intervene may not save it from being “solely automated.” The line between “housekeeping” and “decision” thus likely turns on **impact and discretion**. A trivial algorithmic action that has no appreciable effect on anyone’s rights (say automatically reformatting an application form) is not a decision about a person. But an automated step that **affects whether, how, or by whom a person’s case is examined** – thereby influencing the chances of success – may well fall within Article 22’s ambit. Until further guidance emerges, public authorities using AI in procedural roles should err on the side of providing Article 22 safeguards (notice, the possibility of human review, etc.) for such automated actions if they **materially influence individual outcomes**.

5. Profiling through content analysis: when does AI evaluation of submissions profile the individual?

GDPR Article 4(4) defines **“profiling”** as any automated processing of personal data intended **“to evaluate certain personal aspects relating to a natural person”**, particularly to **“analyse or predict aspects**

concerning that natural person's performance, economic situation, health, personal preferences, interests, reliability, behaviour, location or movements." [17] In most Article 22 cases, profiling is thought of in terms of evaluating *the person* – e.g. their creditworthiness, their job performance, their health risks. However, AI systems might be deployed to evaluate *content* provided by the person, such as the arguments in a consultation submission or the tone of a message. Does this constitute profiling of the individual?

The answer depends on whether the system is effectively assessing personal qualities of the individual through that content. Often, it is. **Sentiment analysis**, for instance, uses a person's statements to infer their emotional state or attitude (e.g. labeling a comment as indicating anger, positivity, fear, etc.). That clearly is an evaluation of personal aspects – the author's emotions or mindset – and thus fits squarely within the GDPR's profiling definition. Similarly, if an algorithm analyzes a person's writing style or substance to judge credibility or to compare it against past behavior (for anomalies), it is drawing inferences about the person's **reliability or behavior pattern**. Indeed, the *SCHUFA* credit scoring was essentially content-agnostic profiling: a numerical value predicting someone's future behavior (likelihood of repaying a loan) was deemed profiling [18]. By analogy, an AI that flags a public consultation submission as an outlier (perhaps because it doesn't fit usual patterns of feedback) is implicitly making a statement about the **author** – e.g. that this person's input is unusually negative or irrelevant compared to others. This too can be seen as evaluating a personal aspect (their viewpoint or the seriousness of their concern).

That said, there is a gray area. Suppose an AI tool purely evaluates the *content's merits*: for example, ranking arguments by their relevance to the consultation question. One could argue the system is not assessing the individual per se, only the document. If no attribute of the author is inferred – if the analysis stops at "this comment is on-topic or off-topic" – is that profiling? The counterpoint is that even labeling a person's contribution as off-topic assigns them a characteristic in context (they *made* an off-topic comment, implying something about their understanding or intentions). GDPR authorities have tended to take a broad view of what counts as profiling, emphasizing context and downstream use. If the result of content analysis is used to treat the person differently (e.g. disregarding their input, or subjecting them to greater scrutiny), it will likely be deemed profiling **because it influences the way the person is evaluated and handled**.

In sum, **most AI-driven evaluation of user-provided content will amount to profiling** if it involves any judgment or inference about the person who provided it. The legal ramifications are that such processing must comply with GDPR's profiling and transparency rules, and if it leads to an automated decision with significant effects, Article 22 kicks in. The uncertainty lies in edge cases – e.g., purely semantic classification with no bearing on the person's qualities – but those are rare in practice. As a practical matter, treating content evaluation algorithms as profiling is the safer course, ensuring data subjects retain rights over these inferred assessments.

6. Satisfying Article 22(2)(b): suitable safeguards for automated administrative decisions

Legal authorization and safeguards: Article 22(2)(b) GDPR creates an exception to the prohibition of solely automated decisions when the decision **"is authorised by Union or Member State law"** that applies to the controller **"and which also lays down suitable measures to safeguard the data subject's rights and freedoms and legitimate interests."** In an administrative context, this means a Member State can, by law, permit public authorities to make decisions about individuals using AI – but **only if that law is**

appropriately framed. Two key requirements emerge: (1) the law must be **clear, precise, and foreseeable** in its terms (a general principle for any legal basis under GDPR [19]), and (2) the law must specify **safeguards** for individuals affected by the automated decision. Simply authorizing "administrative agencies may use automated decision-making" in vague terms would not suffice; the legislation should delineate how the individual's rights are protected in the process [19] [20].

What counts as "suitable measures" is not exhaustively listed in GDPR. Article 22(3) provides a baseline for cases of contract necessity or consent: at minimum, the data subject should have **the right to obtain human intervention, to express their point of view, and to contest the decision** [21]. These rights reflect the spirit of due process – a second chance for the individual to have a human review or to present their side. Although Article 22(3) does not explicitly mandate the same safeguards for decisions made under a statutory authorization (point (b)), many commentators and the logic of the law suggest that **similar safeguards should usually be provided in that context** [22] [23]. Recital 71 GDPR even adds that the person should have the right to **"obtain an explanation of the decision reached after such assessment"**, implying transparency is part of the package of safeguards in automated decision scenarios [24]. In short, a national law that permits automated administrative decisions is expected to build in protections akin to notice, the ability to seek human review, an opportunity to be heard, and avenues to challenge the outcome.

General administrative law vs. data protection-specific safeguards: Many Member States have general administrative procedure acts that guarantee fundamental fairness: rights to be notified of decisions, to receive a reasoned explanation, to appeal to an independent body, etc. [25]. Do these general rights satisfy Article 22(2)(b)'s requirement of "suitable measures," or must the legislation include **extra, AI-specific provisions**? This remains an open question, as neither the CJEU nor the EDPB has given detailed guidance on it. In practice, we see different approaches. Some countries (for example, France in its administrative code, or specific Scandinavian laws for social benefits) have introduced explicit clauses addressing algorithmic decisions – obliging the administration to inform citizens when a decision is algorithmically generated and sometimes to disclose the logic or rule set upon request. These specific measures directly target the novel risks of AI decisions (opacity, lack of human empathy, etc.). Other countries have proceeded on the assumption that existing procedural guarantees (like the right to appeal an unfavorable decision to a tribunal) inherently provide a check on automated decisions, since a person can always have a human judge ultimately review the case.

There is a consensus, however, that **general principles must be respected with even greater vigilance when decisions are automated**. For example, the principle of giving reasons for a decision might require *more* detail if the decision was made by an algorithm – otherwise the individual cannot meaningfully challenge it. If a law authorizes an authority to decide by algorithm, at the very least it should not exempt that authority from the usual duty to explain the decision to the person (in understandable terms). Additionally, Recital 41 GDPR emphasizes that any legal authorisation for automated decisions must be **"clear and precise and its application should be foreseeable"** to those subject to it [19]. A vague empowerment (e.g. "the agency may use computer programs to assist in decision-making") without concrete safeguards would likely be viewed as **insufficient under GDPR** [26]. The *SCHUFA* judgment touched on this issue in passing: it noted that Member State law cannot override GDPR rights or water down the balancing of interests under Article 6 – for instance, a national law can't predetermine that automated processing is always allowed without case-by-case fairness checks [27] [28]. By analogy, a national provision that simply declares automated processing lawful, but fails to set out how individuals can exercise their

rights, might be found not to meet the **"suitable measures"** threshold or even conflict with the GDPR's core principles.

In absence of explicit precedent, a prudent interpretation is that **administrative AI use must embed both traditional administrative safeguards and the specific rights from data protection law**. Existing administrative remedies (like the right to appeal to a court or ombudsman) are indispensable but might not be enough on their own. Data protection demands proactive measures: informing individuals that a decision was automated (GDPR Articles 13–15 require this disclosure [29] [30]), allowing them to request human intervention and present additional input, and guarding against bias or error through technical and organizational measures. Indeed, best practice suggestions (from WP29 and others) include **algorithmic audits, accuracy and non-discrimination testing, and strict data minimisation** as part of the safeguard toolbox [31].

In conclusion, **Member State laws authorising AI in public decision-making should be narrowly tailored**. They must articulate the scope of automation and incorporate safeguards equivalent to those in Article 22(3). Generic administrative law principles – like the right to a fair hearing, reasoned decision, and appeal – provide a safety net, but without explicit alignment to Article 22, there's a risk they may not fully satisfy GDPR requirements. As the use of AI in government expands, we can expect more scrutiny on whether legislation meets the **"clear and precise" standard with "suitable measures"** built in. Until courts clarify this, governments would be wise to err on the side of specificity: expressly writing into law the rights of individuals when an algorithmic tool is used, such as the right to be informed about the AI's role and to a meaningful review by a human. This not only helps comply with Article 22(2)(b) but also upholds citizens' trust in the fairness of automated administrative processes.

Sources:

[^1]: GDPR Article 22(1) (Regulation (EU) 2016/679); EDPB, *Guidelines on Automated individual decision-making and Profiling*, WP251 rev.01, p. 21 (adopted 6 Feb. 2018) [1] [3].

[^2]: Autoriteit Persoonsgegevens (Netherlands DPA), *Meaningful Human Intervention* (Consultation document, March 2025) at 17–18 (examples of factors for meaningful oversight) [5].

[^3]: CJEU, *Case C-634/21, SCHUFA Holding AG* (judgment of 7 Dec. 2023) paras 45–50 [6] [10].

[^4]: Ruth Boardman & Nils Lölfing, IAPP Analysis of CJEU *SCHUFA* ruling (19 Dec. 2023) (noting CJEU's rejection of the "preparatory act" argument) [8] [9].

[^5]: AEPD (Spain), Blog: *Evaluating human intervention in automated decisions* (4 March 2024) (summary of CJEU *SCHUFA* and its implications) [11].

[^6]: Chantal Prudhomme et al., *Administrative Law and AI Decision-Making* in *EU Law Journal* (2024) (discussing national case law post-*SCHUFA*, including Amazon warehouse case) [12].

[^7]: Aarhus Convention 1998, Art. 6(8) (public participation) ("each Party shall ensure that due account is taken of the outcome of the public participation") [15]; see also EU Directive 2011/92 (EIA Directive) Art. 8 (requiring that results of consultations be considered in decision).

[^8]: Article 29 Data Protection Working Party, *Guidelines on Automated Decision-making* (WP251, 2018) (on multi-stage processing and locating the "decision") [16].

[^9]: ICO (UK), *Guidance on AI and Data Protection* (2020), Part 3 (human involvement must be after the AI and affect the outcome to be meaningful).

[^10]: GDPR Article 4(4) (definition of profiling) [17].

[^11]: CJEU, *SCHUFA* (2023) (confirming that generating a probability score about an individual is profiling and an automated decision) [18].

[^12]: Regulation (EU) 2016/679, Recital 41 (law authorising automated decisions must be clear, precise and foreseeable) [19].

[^13]: CJEU, *SCHUFA* (2023), paras 74–80 (national law permitting automated scoring must still comply with GDPR, and cannot predetermine legitimate interest balancing) [27] [28].

[^14]: Bygrave, Lee A., *"Minding the Machine v2.0"*, in *Computer Law & Security Review* 2020, p. 105 (suggesting Article 22(2)(b) laws should include safeguards akin to Art 22(3)) [22] [23].

[^15]: *Proposal for an EU Regulation on AI Liability* (2022) (noting the importance of transparency and human oversight in public-sector AI deployments).

[^16]: Weitzenboeck, Emily M., *"Fully Automated Administrative Decisions in Norway"*, *Adm.Sci.* 11:149 (2021) (comparative analysis of safeguards in Norwegian law) [21] [25].

[^17]: GDPR Recital 71 (data subjects should have the right to an explanation of automated decisions and other safeguards) [24].

[^18]: Article 29 Working Party, WP251 (2018) (additional recommended safeguards: algorithmic auditing, bias testing, etc.) [31].

---

[1] [2] [3] [4] [5] [11] Evaluating human intervention in automated decisions | AEPD
https://www.aepd.es/en/press-and-communication/blog/evaluating-human-intervention-in-automated-decisions

[6] [7] [10] [13] [14] [17] [18] [29] [30] CURIA - Documents
https://curia.europa.eu/juris/document/document.jsf?docid=280426&doclang=en

[8] [9] [27] [28] Key takeaways from the CJEU's recent automated decision-making rulings | IAPP
https://iapp.org/news/a/key-takeaways-from-the-cjeus-recent-automated-decision-making-rulings

[12] [16] Scores as Decisions? Article 22 GDPR and the Judgment of the CJEU in SCHUFA Holding (Scoring) in the Labour Context | Industrial Law Journal | Oxford Academic
https://academic.oup.com/ilj/article/53/4/840/7745471

15 Chapter 10 The Human Rights Dimension in: Sustainability ... - Brill
https://brill.com/edcollchap-oa/book/9789004509382/BP000021.xml?language=en&srsltid=AfmBOoqN5XGu7t5DV8cb6TdXIwuxFEdHc0Km8qLEB_nnm8XyjxcXOOla

19 20 21 22 23 24 25 26 31 Simplification of Administrative Procedures through Fully Automated Decision-Making: The Case of Norway
https://www.mdpi.com/2076-3387/11/4/149